

## Contents

<b>Monitoring HPE GreenLake Servers running GPU using Grafana and Prometheus</b>	<b>2</b>
Overview . . . . .	2
Kubernetes and Helm Setup . . . . .	2

# Monitoring HPE GreenLake Servers running GPU using Grafana and Prometheus

## Overview

HPE GreenLake provides a cloud-native platform for managing and monitoring infrastructure with built-in tools and dashboards. While GreenLake offers comprehensive native monitoring capabilities, organizations can also leverage the GreenLake API to integrate with popular open-source tools like Grafana and Prometheus. This approach enables teams to consolidate monitoring data across hybrid environments, utilize existing observability workflows, and create customized dashboards tailored to specific operational needs.

## Kubernetes and Helm Setup

### Kubernetes cluster setup

This demonstration environment utilizes a high-availability Kubernetes cluster consisting of three control plane nodes and two worker nodes,

```
wsl=> k get node -o wide
```

NAME	EXTERNAL-IP	OS-IMAGE	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	
				KERNEL-VERSION		CONTAINER-RUNTIME		
c2-cp-01.hst.enablement.local		Ubuntu 22.04.5 LTS	Ready	control-plane	80d	v1.32.5	10.16.160.51	<
none>				5.15.0-144-generic		containerd://2.0.5		
c2-cp-02.hst.enablement.local		Ubuntu 22.04.5 LTS	Ready	control-plane	80d	v1.32.5	10.16.160.52	<
none>				5.15.0-144-generic		containerd://2.0.5		
c2-cp-03.hst.enablement.local		Ubuntu 22.04.5 LTS	Ready	control-plane	80d	v1.32.5	10.16.160.53	<
none>				5.15.0-144-generic		containerd://2.0.5		
c2-worker-01.hst.enablement.local		Ubuntu 22.04.5 LTS	Ready	<none>	80d	v1.32.5	10.16.160.54	<
none>				5.15.0-144-generic		containerd://2.0.5		
c2-worker-02.hst.enablement.local		Ubuntu 22.04.5 LTS	Ready	<none>	80d	v1.32.5	10.16.160.55	<
none>				5.15.0-144-generic		containerd://2.0.5		

### Kubernetes namespace setup

The cluster is equipped with the [gpu-operator](#) namespace for NVIDIA GPU management and the [monitoring](#) namespace hosting the Prometheus stack, with external access enabled via NodePort services.

```
wsl=> kubectl get ns | grep -vE '^(kube-|default)'
```

NAME	STATUS	AGE
gpu-operator	Active	80d
monitoring	Active	56d

```
wsl=> k get svc -n gpu-operator
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

gpu-operator	ClusterIP	10.233.44.80	<none>	8080/TCP	78d
nvidia-dcgm-exporter	ClusterIP	10.233.15.59	<none>	9400/TCP	78d
wsl=> k get svc --field-selector spec.type=NodePort -n monitoring					
NAME	AGE	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kube-prometheus-stack-grafana	56d	NodePort	10.233.22.241	<none>	80:30080/TCP
kube-prometheus-stack-prometheus	56d	NodePort	10.233.8.106	<none>	9090:30090/TCP

## Helm chart installation

The environment uses Helm to manage two key components: the NVIDIA GPU Operator for GPU resource management and the Kube Prometheus Stack for monitoring and observability.

wsl=> helm list -A					
NAME	NAMESPACE	REVISION	UPDATED		APP
VERSION		STATUS	CHART		
gpu-operator-1753140595	gpu-operator	4	2025-08-14 19:20:42.329819669	-0700	
MST deployed	gpu-operator-v25.3.2		v25.3.2		
kube-prometheus-stack	monitoring	5	2025-08-15 13:06:31.169338089	-0700	
MST deployed	kube-prometheus-stack-76.3.				
0	v0.84.1				

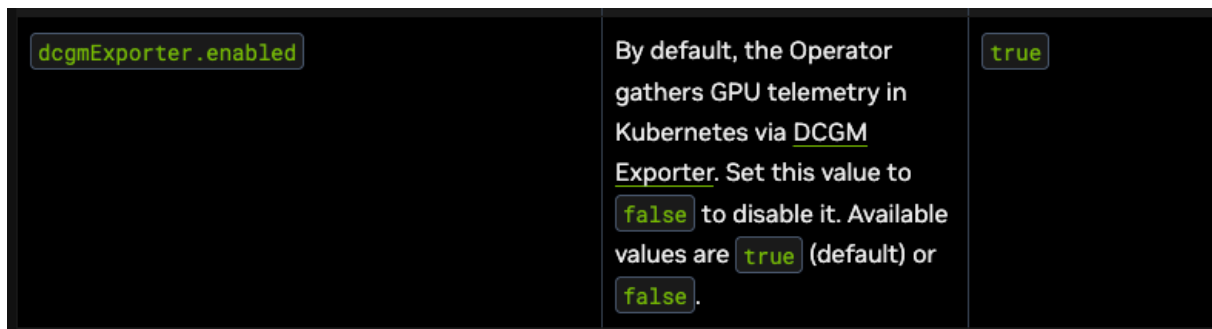
## GPU Operator chart customization

The NVIDIA GPU Operator Helm chart deploys a DCGM (Data Center GPU Manager) exporter by default, but there are important nuances:

- The DCGM exporter Pod will be created automatically when the operator detects a node with an NVIDIA GPU and the dcgm-exporter component is enabled in its values.

wsl=> k -n gpu-operator get pods -o wide   grep dcgm					
nvidia-dcgm-exporter-gkg6d			1/1	Running	0
56d	10.233.117.209	c2-worker-01.hst.enablement.local	<none>		<
none>					
nvidia-dcgm-exporter-r2np6			1/1	Running	0
56d	10.233.114.15	c2-worker-02.hst.enablement.local	<none>		<
none>					

- In the stock gpu-operator Helm chart from NVIDIA's repo, the DCGM exporter is enabled by default (`dcgmExporter.enabled: true`). See the [NVIDIA GPU Operator Documentation](#).



**Figure 1:** GPU Operator Configuration

However:

1. ServiceMonitor is not enabled by default.
  - This means Prometheus won't automatically scrape the DCGM exporter unless you either:
    - Enable the ServiceMonitor (`dcmExporter.serviceMonitor.enabled: true`), or
    - Manually define a scrape config in Prometheus.

The gpu-operator is configured with custom values to enable Prometheus integration. The DCGM exporter runs as a ClusterIP service with ServiceMonitor enabled for automatic metrics discovery by Prometheus.

```
wsl=> helm get values gpu-operator-1753140595 -n gpu-operator
USER-SUPPLIED VALUES:
dcmExporter:
  service:
    type: ClusterIP
  serviceMonitor:
    enabled: true
```

## GPU utilization simulation

To simulate GPU load and verify monitoring functionality, we deployed a test pod running the gpu-burn utility. This tool performs intensive GPU computations, allowing us to observe GPU utilization metrics in our monitoring dashboards.

The following YAML manifest creates a pod that clones the gpu-burn repository, compiles it, and runs continuous GPU stress testing:

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-burn
```

```
spec:
  containers:
    - name: gpu-burn
      image: nvidia/cuda:12.2.0-devel-ubuntu22.04
      command: ["/bin/bash", "-c"]
      args:
        - |
          apt update && apt install -y git build-essential && \
          git clone https://github.com/wilicc/gpu-burn.git && \
          cd gpu-burn && make && ./gpu_burn 999999
      resources:
        limits:
          nvidia.com/gpu: 1
      restartPolicy: Never
```

**Key configuration details:** - **Base image:** `nvidia/cuda:12.2.0-devel-ubuntu22.04` provides the CUDA development environment

- **GPU allocation:** `nvidia.com/gpu: 1` requests a single GPU from the cluster
- **Runtime:** `gpu_burn 999999` runs for approximately 277 hours (effectively continuous)
- **Restart policy:** `Never` ensures the pod completes its run without automatic restarts

Deploy the pod using:

```
kubectl apply -f gpu-burn.yaml
```