

Contents

Monitoring HPE GreenLake Servers running GPU using Grafana and Prometheus	2
Overview	2
Kubernetes and Helm Setup	2

Monitoring HPE GreenLake Servers running GPU using Grafana and Prometheus

Overview

HPE GreenLake provides a cloud-native platform for managing and monitoring infrastructure with built-in tools and dashboards. While GreenLake offers comprehensive native monitoring capabilities, organizations can also leverage the GreenLake API to integrate with popular open-source tools like Grafana and Prometheus. This approach enables teams to consolidate monitoring data across hybrid environments, utilize existing observability workflows, and create customized dashboards tailored to specific operational needs.

Kubernetes and Helm Setup

Environment Verification

Before proceeding with the monitoring setup, verify that your Kubernetes cluster has the necessary components installed. The following shows a working environment with the GPU Operator and Prometheus monitoring stack deployed:

Services running in the `gpu-operator` namespace: - `gpu-operator`: Core service for GPU management (ClusterIP: 10.233.44.80:8080) - `nvidia-dcgm-exporter`: DCGM metrics exporter for Prometheus integration (ClusterIP: 10.233.15.59:9400)

Helm releases: - `gpu-operator-1753140595` (v25.3.2) in the `gpu-operator` namespace - `kube-prometheus-stack` (76.3.0) in the `monitoring` namespace

You can verify your setup using the following commands:

```
wsl=> k get svc -n gpu-operator
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
gpu-operator	ClusterIP	10.233.44.80	<none>	8080/TCP	78d
nvidia-dcgm-exporter	ClusterIP	10.233.15.59	<none>	9400/TCP	78d


```
wsl=> helm list -A
```

NAME	NAMESPACE	REVISION STATUS	UPDATED CHART	APP
VERSION				
gpu-operator-1753140595	gpu-operator	4	2025-08-14 19:20:42.329819669	-0700
MST deployed	gpu-operator-v25.3.2		v25.3.2	
kube-prometheus-stack	monitoring	5	2025-08-15 13:06:31.169338089	-0700
MST deployed	kube-prometheus-stack-76.3.0		v0.84.1	

```
hjma@HSTHJMA02:~
```

External Access Configuration

The cluster has been configured with NodePort services to enable external access to Grafana and Prometheus:

```
wsl=> k get svc --field-selector spec.type=NodePort
```

NAME	AGE	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kube-prometheus-stack-grafana	55d	NodePort	10.233.22.241	<none>	80:30080/TCP
kube-prometheus-stack-prometheus	55d	NodePort	10.233.8.106	<none>	9090:30090/TCP

GPU utilization simulation

To simulate GPU load and verify monitoring functionality, we deployed a test pod running the gpu-burn utility. This tool performs intensive GPU computations, allowing us to observe GPU utilization metrics in our monitoring dashboards.

The following YAML manifest creates a pod that clones the gpu-burn repository, compiles it, and runs continuous GPU stress testing:

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-burn
spec:
  containers:
    - name: gpu-burn
      image: nvidia/cuda:12.2.0-devel-ubuntu22.04
      command: ["/bin/bash", "-c"]
      args:
        - |
          apt update && apt install -y git build-essential && \
          git clone https://github.com/wilicc/gpu-burn.git && \
          cd gpu-burn && make && ./gpu_burn 999999
      resources:
        limits:
          nvidia.com/gpu: 1
      restartPolicy: Never
```

Key configuration details: - **Base image:** `nvidia/cuda:12.2.0-devel-ubuntu22.04` provides the CUDA development environment - **GPU allocation:** `nvidia.com/gpu: 1` requests a single GPU from the cluster - **Runtime:** `gpu_burn 999999` runs for approximately 277 hours (effectively continuous) - **Restart policy:** `Never` ensures the pod completes its run without automatic restarts

Deploy the pod using:

```
kubectl apply -f gpu-burn.yaml
```